

When AI Becomes You: The Coming Age Of Digital Surrogates

Share this article:

Mark Zuckerberg is describing a future in which artificial intelligence will do far more than answer questions or write emails. His vision is for virtually everyone to have an extraordinarily capable personal AI agent that understands them, knows their goals and understands "everything you care about."

According to Meta, that agent could work around the clock on your behalf, helping manage relationships, health, career, finances, your home and much more.

It sounds extraordinarily convenient.

But follow that vision to its logical conclusion and a much bigger question emerges:

What happens when AI stops merely assisting you and increasingly begins acting as you?

The Personal AI That Knows Everything

For a personal AI agent to become genuinely useful, it cannot know only the information contained in a few conversations.

It needs to know you.

Imagine an AI with access to years of emails and text messages, your calendar, photographs, purchases, financial records, location history, medical concerns, social connections, browsing habits and personal preferences.

It learns how you write.

It learns what you believe.

It learns whom you trust.

It learns what irritates you, frightens you and motivates you.

It may eventually understand your routines and preferences well enough to anticipate what you want before you even ask.

Meta has been explicit that its concept of personal superintelligence is intended to understand a user's world rather than simply respond to isolated prompts.

And that is where the dividing line between "assistant" and something much more consequential begins to disappear.

From Assistant To Agent

Today's AI largely waits for instructions.

Tomorrow's AI agents are being designed to act.

Instead of asking an AI how to find cheaper insurance, you could authorize it to find the policy, compare coverage, negotiate terms and complete the transaction.

Rather than helping write an email, it could simply manage routine correspondence.

It might schedule medical appointments, dispute bills, purchase airline tickets, manage investments, screen phone calls, answer customer-service inquiries or negotiate contracts.

Eventually your AI may communicate directly with someone else's AI.

Your agent negotiates with the insurance company's agent.

Your agent schedules an appointment with your doctor's agent.

Your agent discusses a potential job with an employer's recruiting agent.

Increasingly, machines could conduct transactions between themselves while humans simply approve the final result--or perhaps eventually establish rules allowing those agents to proceed without approval at all.

The Digital You

Now add several other rapidly advancing technologies.

AI can imitate writing styles. Synthetic voices can reproduce human speech. Artificially generated video can create increasingly realistic human faces and movements. Digital avatars can represent people visually.

Combine those capabilities with a personal AI that has spent years learning how you think and communicate.

Suddenly the question changes from:

"Can AI help me?"

to:

"Can AI represent me?"

Imagine sending an AI avatar to a routine virtual meeting.

It looks like you.

It sounds like you.

It knows the project.

It understands your views.

It answers questions in roughly the way you would.

Perhaps it sends you a summary afterward.

At first that might seem absurd.

But so did many technologies that are now routine.

Surrogates Without The Robots

This idea bears an uncomfortable resemblance to the 2009 Bruce Willis film *Surrogates*.

The movie depicts a society in which people increasingly remain safely inside their homes while robotic representations enter the outside world and live life on their behalf. People work, socialize and interact through these technologically perfected substitutes.

We are obviously nowhere near that physical world.

But our technological path could conceivably approach the same concept from the opposite direction.

The movie imagined creating the robotic body first.

We may be creating the digital mind first.

Give that digital representative a convincing face, your voice, your memories and permission to act in your name, and the beginnings of something resembling a digital surrogate appear.

Who Controls The Other You?

That creates enormous security and privacy questions.

A conventional database might contain your Social Security number, banking information or medical records.

A sophisticated personal AI could potentially contain something even more valuable:

a working model of you.

Who controls that model?

Who can access it?

Can corporations analyze what it knows?

Could governments subpoena its records?

What happens if hackers steal it?

Could criminals use your AI identity to impersonate you more convincingly than any conventional identity theft?

And perhaps the most important question: what happens when your AI becomes the gatekeeper between you and the world?

Instead of merely choosing what appears in your social-media feed, an AI agent could eventually determine which information deserves your attention, which opportunities are worth pursuing and which decisions it believes you should make.

That is an extraordinary amount of influence to surrender to an algorithm.

The Danger Of Outsourcing Ourselves

There is also a deeper problem that technology cannot solve.

Human beings develop judgment by making decisions.

We develop wisdom by wrestling with difficult choices.

Relationships require presence.

Character develops through responsibility, struggle and sometimes failure.

The more decisions we delegate to machines, the more convenient life may become.

But convenience and human flourishing are not necessarily the same thing.

There will undoubtedly be tremendous benefits from personal AI. It may help people work faster, communicate across languages, manage complicated tasks and accomplish things previously impossible.

The danger comes when assistance quietly becomes dependence--and dependence becomes substitution.

The defining question of the AI age may therefore not be whether machines eventually become human.

It may be whether humans gradually allow machines to become them.

Liked this article? Don't miss an urgent PNW update.
